

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.keras import layers, models
4 from tensorflow.keras.datasets import mnist
5 import matplotlib.pyplot as plt
6 from tensorflow.keras.callbacks import ModelCheckpoint
7 import os
8 from skimage.metrics import mean_squared_error, peak_signal_noise_ratio, structural_similarity
```

## ✓ Loading and Transforming

```
1 (x_train, _), (x_test, _) = mnist.load_data() # Loads the dataset and ignore the labels
2 x_train = x_train.astype("float32") / 255.0 # Normalizes the training images to the range [0, 1]
3 x_test = x_test.astype("float32") / 255.0 # Normalizes the test images to the range [0, 1]
```

📄 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 ————— 0s 0us/step

```
1 x_train.shape
```

📄 (60000, 28, 28)

```
1 x_test.shape
```

📄 (10000, 28, 28)

## ✓ Adding Noise at Random Locations

```
1 def add_noise_random_locations(images, noise_level=0.5, noise_range=[0, 255]):
2     """
3     Add noise at random locations within the image.
4
5     Parameters:
6     images (numpy.ndarray): Input images.
7     noise_level (float): The proportion of pixels to add noise to.
8     noise_range (list): The range of noise values to add.
9
10    Returns:
11    numpy.ndarray: Noisy images with pixel values clipped to [0, 1].
12    """
13    noisy_images = np.copy(images)
14    num_pixels = images.shape[1] * images.shape[2] # 28 * 28 = 784
15    num_noisy_pixels = int(noise_level * num_pixels) # 50% of 784 = 392
16
17    for i in range(images.shape[0]):
18        # Generates random indices to add noise
19        noisy_indices = np.random.choice(num_pixels, num_noisy_pixels, replace=False)
20        noisy_indices = np.unravel_index(noisy_indices, images.shape[1:])
21
22        # Adds noise within the specified range
23        noise = np.random.uniform(noise_range[0], noise_range[1], size=num_noisy_pixels) / 255.0
24        noisy_images[i][noisy_indices] = noise
25
26    return np.clip(noisy_images, 0., 1.) # np.clip(array, min_value, max_value)
```

```
1 # Generate noisy training and test data
2 x_train_noisy = add_noise_random_locations(x_train)
3 x_test_noisy = add_noise_random_locations(x_test)
```

## ✓ Convolutional Autoencoder Model

```
1 input_img = layers.Input(shape=(28, 28, 1))
```

```
1 # Encoder
2 x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
3 x = layers.MaxPooling2D((2, 2), padding='same')(x)
4 x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
5 encoded = layers.MaxPooling2D((2, 2), padding='same')(x)
```

```
1 # Decoder
2 x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(encoded)
3 x = layers.UpSampling2D((2, 2))(x)
4 x = layers.Conv2D(32, (3, 3), activation='relu', padding='same')(x)
5 x = layers.UpSampling2D((2, 2))(x)
6 decoded = layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)
```

```
1 # Creates the autoencoder model
2 autoencoder = models.Model(input_img, decoded)
3 autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

```
1 print("Autoencoder Model Summary:")
2 autoencoder.summary()
```



Autoencoder Model Summary:  
Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 28, 28, 1)	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_2 (Conv2D)	(None, 7, 7, 32)	9,248
up_sampling2d (UpSampling2D)	(None, 14, 14, 32)	0
conv2d_3 (Conv2D)	(None, 14, 14, 32)	9,248
up_sampling2d_1 (UpSampling2D)	(None, 28, 28, 32)	0
conv2d_4 (Conv2D)	(None, 28, 28, 1)	289

Total params: 28,353 (110.75 KB)  
Trainable params: 28,353 (110.75 KB)

## Training

```
1 # Reshapes the data to include the channel dimension
2 x_train = np.reshape(x_train, (-1, 28, 28, 1))
3 x_test = np.reshape(x_test, (-1, 28, 28, 1))
4 x_train_noisy = np.reshape(x_train_noisy, (-1, 28, 28, 1))
5 x_test_noisy = np.reshape(x_test_noisy, (-1, 28, 28, 1))
```

```
1 path = '/content/drive/MyDrive/ProgressSoft /Image Noise'
```

```
1 full_path = os.path.join(path, "training_1")
```

```
1 os.makedirs(full_path, exist_ok=True)
```

```
1 # Define the checkpoint callback
2 checkpoint_path = os.path.join(full_path, "cp.weights.h5")
3 checkpoint_dir = os.path.dirname(checkpoint_path)
4 checkpoint_callback = ModelCheckpoint(
5     filepath=checkpoint_path,
6     save_weights_only=True,
7     save_freq='epoch',
8     verbose=1
9 )
```

```
1 # Train the model with the checkpoint callback
2 autoencoder.fit(
3     x_train_noisy, x_train,
4     epochs=40,
5     batch_size=256,
6     validation_data=(x_test_noisy, x_test),
7     callbacks=[checkpoint_callback]
8 )
```



235/235 ————— 141s 527ms/step - loss: 0.0929 - val\_loss: 0.0920  
Epoch 28/40  
235/235 ————— 0s 507ms/step - loss: 0.0929  
Epoch 28: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 143s 531ms/step - loss: 0.0929 - val\_loss: 0.0918  
Epoch 29/40  
235/235 ————— 0s 512ms/step - loss: 0.0924  
Epoch 29: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 143s 535ms/step - loss: 0.0924 - val\_loss: 0.0918  
Epoch 30/40  
235/235 ————— 0s 512ms/step - loss: 0.0922  
Epoch 30: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 142s 535ms/step - loss: 0.0922 - val\_loss: 0.0914  
Epoch 31/40  
235/235 ————— 0s 510ms/step - loss: 0.0920  
Epoch 31: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 142s 534ms/step - loss: 0.0920 - val\_loss: 0.0910  
Epoch 32/40  
235/235 ————— 0s 500ms/step - loss: 0.0918  
Epoch 32: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 140s 526ms/step - loss: 0.0918 - val\_loss: 0.0912  
Epoch 33/40  
235/235 ————— 0s 508ms/step - loss: 0.0915  
Epoch 33: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 142s 526ms/step - loss: 0.0915 - val\_loss: 0.0908  
Epoch 34/40  
235/235 ————— 0s 513ms/step - loss: 0.0913  
Epoch 34: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 143s 530ms/step - loss: 0.0913 - val\_loss: 0.0906  
Epoch 35/40  
235/235 ————— 0s 505ms/step - loss: 0.0912  
Epoch 35: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 124s 529ms/step - loss: 0.0912 - val\_loss: 0.0905  
Epoch 36/40  
235/235 ————— 0s 497ms/step - loss: 0.0909  
Epoch 36: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 139s 518ms/step - loss: 0.0909 - val\_loss: 0.0903  
Epoch 37/40  
235/235 ————— 0s 505ms/step - loss: 0.0906  
Epoch 37: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 123s 521ms/step - loss: 0.0906 - val\_loss: 0.0901  
Epoch 38/40  
235/235 ————— 0s 500ms/step - loss: 0.0908  
Epoch 38: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 142s 523ms/step - loss: 0.0908 - val\_loss: 0.0898  
Epoch 39/40  
235/235 ————— 0s 501ms/step - loss: 0.0904  
Epoch 39: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 124s 528ms/step - loss: 0.0904 - val\_loss: 0.0899  
Epoch 40/40  
235/235 ————— 0s 513ms/step - loss: 0.0905  
Epoch 40: saving model to /content/drive/MyDrive/ProgressSoft /Image Noise/training\_1/cp.weights.h5  
235/235 ————— 142s 530ms/step - loss: 0.0905 - val\_loss: 0.0895  
<keras.src.callbacks.history.History at 0x7d036e9a9cd0>

```
1 # Save the entire model (architecture + weights + optimizer state)
2 autoencoder.save("denoising_autoencoder_final.keras") # Recommended .keras format
3 # OR
4 autoencoder.save("denoising_autoencoder_final.h5")      # Legacy HDF5 format
5 print("Final model saved.")
```



WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. Please use the Keras 2.x format instead.



## ✓ Predictions

```
1 # Load the saved model
2 loaded_model = tf.keras.models.load_model("/content/drive/MyDrive/ProgressSoft /Image Noise/Models/denoising_autoencoder_final.keras")
```

```
1 denoised_images = loaded_model.predict(x_test_noisy)
```



313/313 ————— 8s 25ms/step

## ✓ Evaluation Metrics

```
1 # Reshape the data
2 x_test = x_test.reshape(-1, 28, 28) # Original images
3 denoised_images = denoised_images.reshape(-1, 28, 28) # Denoised images
4
5 # Initializes lists to store metric values
6 mse_values = []
```

```
7 psnr_values = []
8 ssim_values = []
9
10 # Compute metrics for each image in the test set
11 for i in range(len(x_test)):
12     original_image = x_test[i]
13     denoised_image = denoised_images[i]
14
15     # Compute MSE
16     mse = mean_squared_error(original_image, denoised_image)
17     mse_values.append(mse)
18
19     # Compute PSNR
20     psnr = peak_signal_noise_ratio(original_image, denoised_image, data_range=1.0)
21     psnr_values.append(psnr)
22
23     # Compute SSIM
24     ssim = structural_similarity(original_image, denoised_image, data_range=1.0)
25     ssim_values.append(ssim)
26
27 # Compute average metrics
28 avg_mse = np.mean(mse_values)
29 avg_psnr = np.mean(psnr_values)
30 avg_ssim = np.mean(ssim_values)
31
32 print(f"Average MSE: {avg_mse:.4f}")
33 print(f"Average PSNR: {avg_psnr:.4f}")
34 print(f"Average SSIM: {avg_ssim:.4f}")
```

↔ Average MSE: 0.0094  
Average PSNR: 20.5885  
Average SSIM: 0.8896

## Visualization of results

```
1 n = 10 # Number of images to display
2 plt.figure(figsize=(20, 6))
3
4 # Original Images
5 for i in range(n):
6     # Display original images
7     ax = plt.subplot(3, n, i + 1) # Create a subplot for the original image
8     plt.imshow(x_test[i].reshape(28, 28), cmap="gray") # Display original image
9     plt.title("Original Image")
10    plt.axis("off")
11
12    # Display noisy images
13    ax = plt.subplot(3, n, i + 1 + n)
14    plt.imshow(x_test_noisy[i].reshape(28, 28), cmap="gray")
15    plt.title("Noisy Image")
16    plt.axis("off")
17
18    # Display denoised images
19    ax = plt.subplot(3, n, i + 1 + 2 * n)
20    plt.imshow(denoised_images[i].reshape(28, 28), cmap="gray")
21    plt.title("Denoised Image")
22    plt.axis("off")
23
24 plt.tight_layout()
25 plt.show()
```



